

Chapter One

This is the first paragraph on the first page. It exists to prove that ordinary body text extracts correctly, in reading order, with a normal font.

This is a second paragraph, on the same page, so paragraph-boundary detection has something to work with.

This whole line is italic, to prove font-based style detection has a real signal to read.

This paragraph is deliberately long so that it overflows across the page boundary without any page break command, forcing the layout engine to split it mid sentence exactly where a real book's justified body text would split naturally, which is exactly the case the page break classifier needs to get right: recognizing that nothing deliberate

happened here, that the words before the boundary and the words after the boundary are really one continuous thought that a human reader would never notice was split across two different pages at all, and that treating this boundary as if it were a meaningful new section would be exactly the kind of mistake the whole point of this project is to avoid making silently, so the paragraph keeps going a while longer still, well past where the page is going to run out of room, on purpose.

This is a short paragraph right after the long one, still on whichever page the long one ends up finishing on, with nothing special about the boundary between them either.

Chapter Two

This paragraph opens a deliberate new chapter, with a heading immediately above it and a fresh page underneath it — the clean, unambiguous case a page-break classifier should get right with high confidence.

A Short Poem

*Line one of the poem,
line two of the poem,
centered and italic,
so verse detection has a fixture.*